


Baseline Runtime Breakdown (50 steps, 50 C4 prompts, unoptimized)



GPT-2 backbone accounts for 77.2% of total inference time